



# Kafka 101/Event Driven Architecture

Amanda Gilbert



# Event Driven Systems

Event Driven Systems are **asynchronous**, allowing **decoupled** services to process events in a scalable, flexible way and to **replay** data when needed.



# STATE

Customer's current email address is jane.doe@gmail.com

Customer's current balance is \$450

There are currently 112 size medium t-shirts in inventory

The patient's temperature is 98.6 degrees fahrenheit



# EVENT

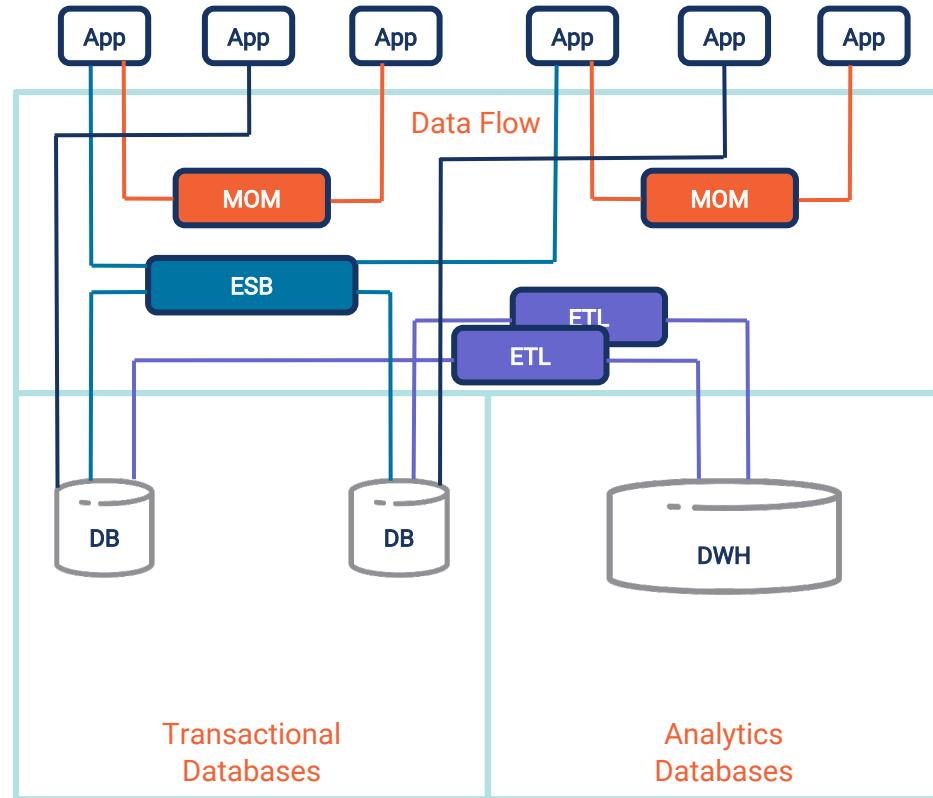
Customer changes email address to jane.doe2@gmail.com

Customer spends \$16 dollars at Target.com

One medium size t-shirt is bought by an online customer

The patient's temperature raises by .8 degrees

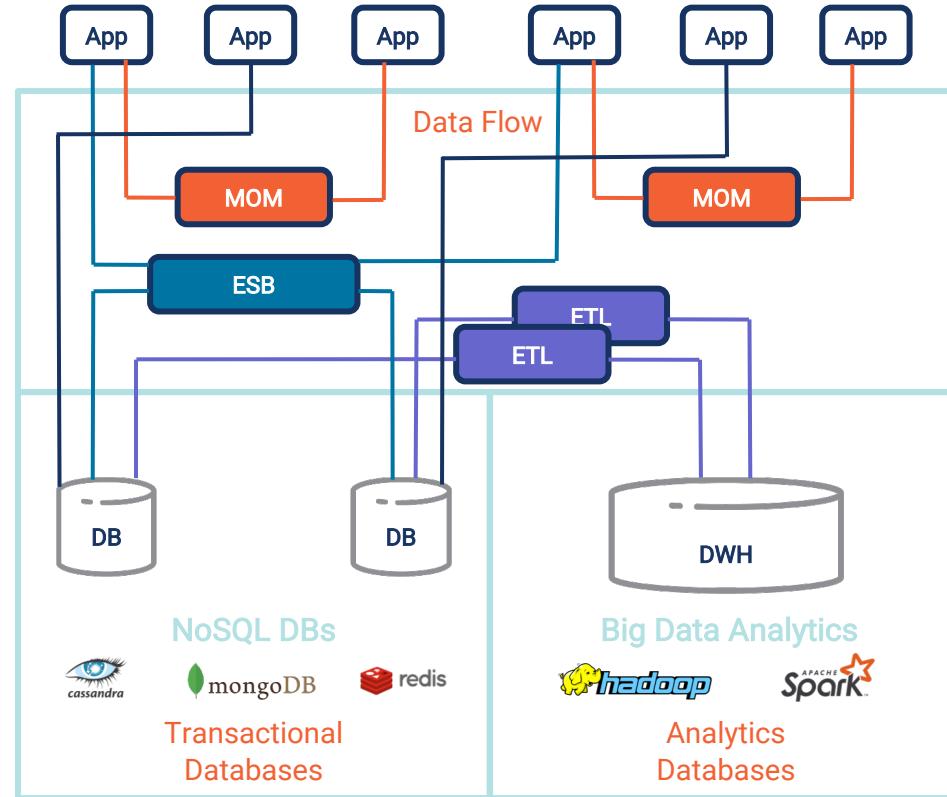
# Legacy Data Infrastructure Solutions Have Architectural Flaws



These solutions can be

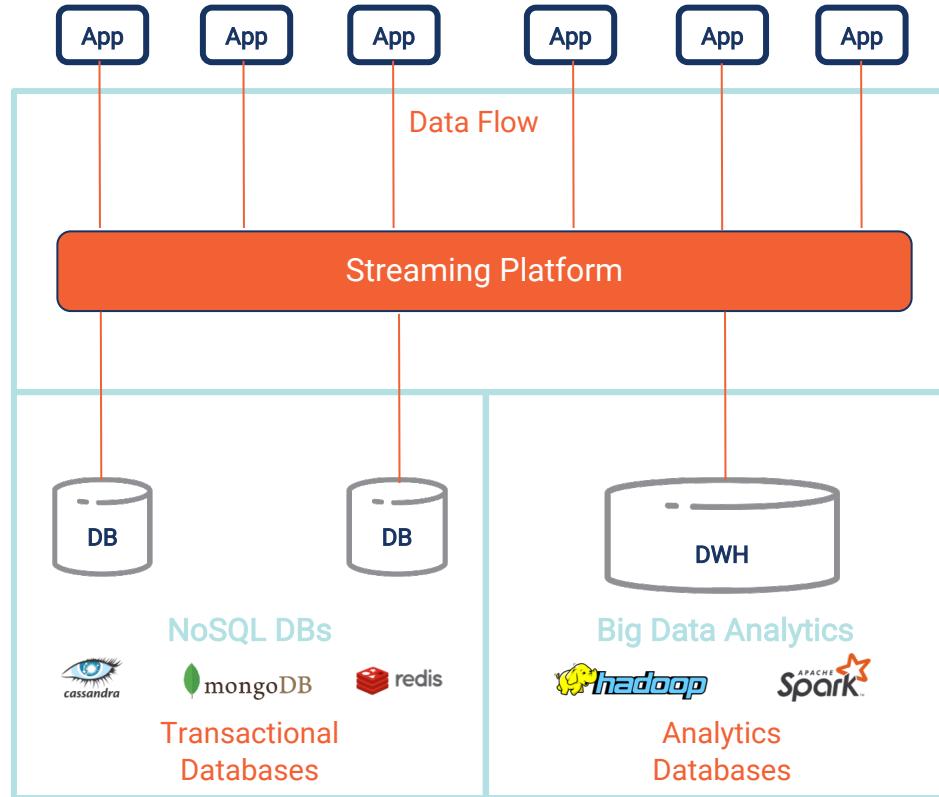
- Batch-oriented, instead of event-oriented in real time
- Complex to scale at high throughput
- Connected point-to-point, instead of publish / subscribe
- Lacking data persistence and retention
- Incapable of in-flight message processing

# Modern Architectures are Adapting to New Data Requirements



But how do we **revolutionize data flow** in a world of exploding, distributed and ever changing data?

# The Solution is a Streaming Platform for Real-Time Data Processing



A Streaming Platform provides a **single source of truth** about your data to everyone in your organization



# **Event Driven Systems pt 2 - EDA Models**



Event Notification

Event-Carried State  
Transfer

Event Sourcing

CQRS



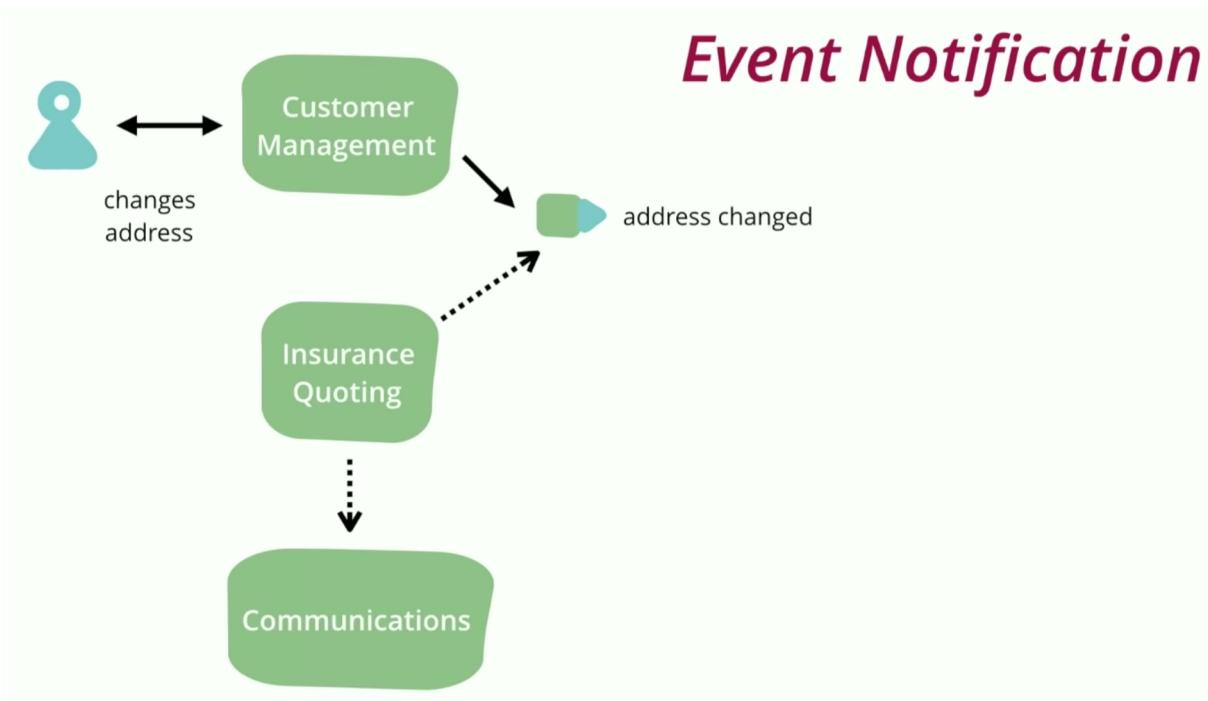
# Event Notification

- Systems notify other systems of events that happen
- Loosely coupled systems, “Fire and forget”
- Usually just metadata transferred

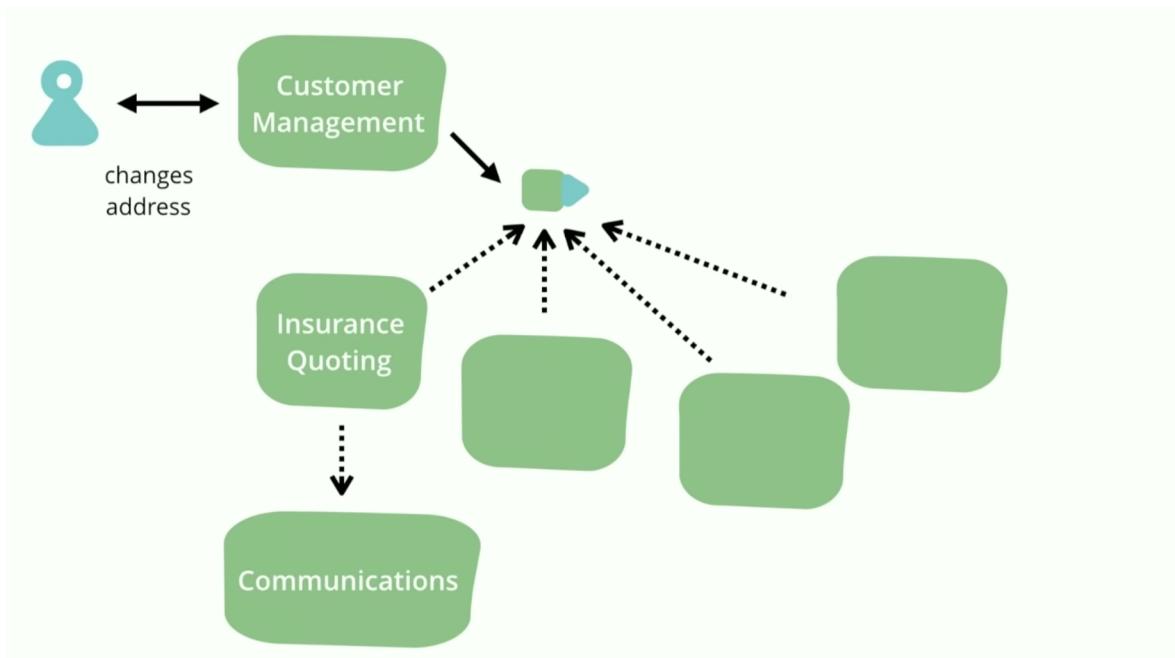
# Use Case: Change customer address resolving dependencies



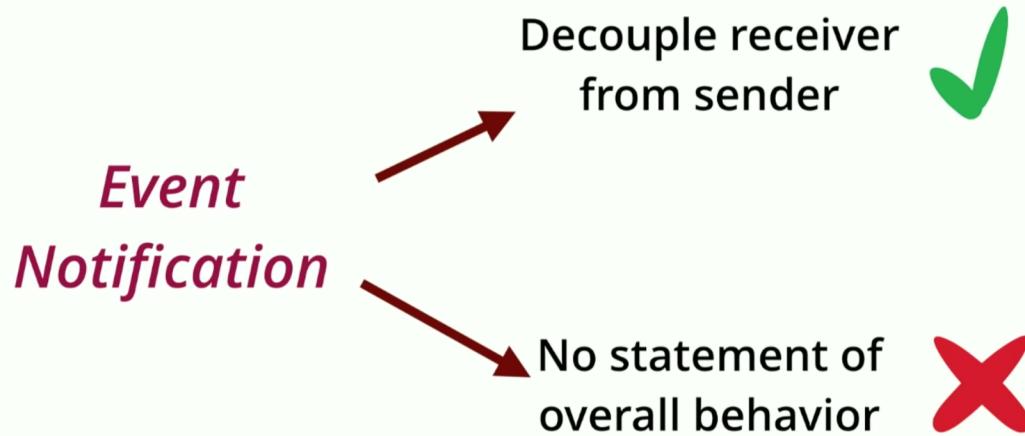
## *Event Notification*



# Use Case: Change customer address, Events or Commands



# Notification Events

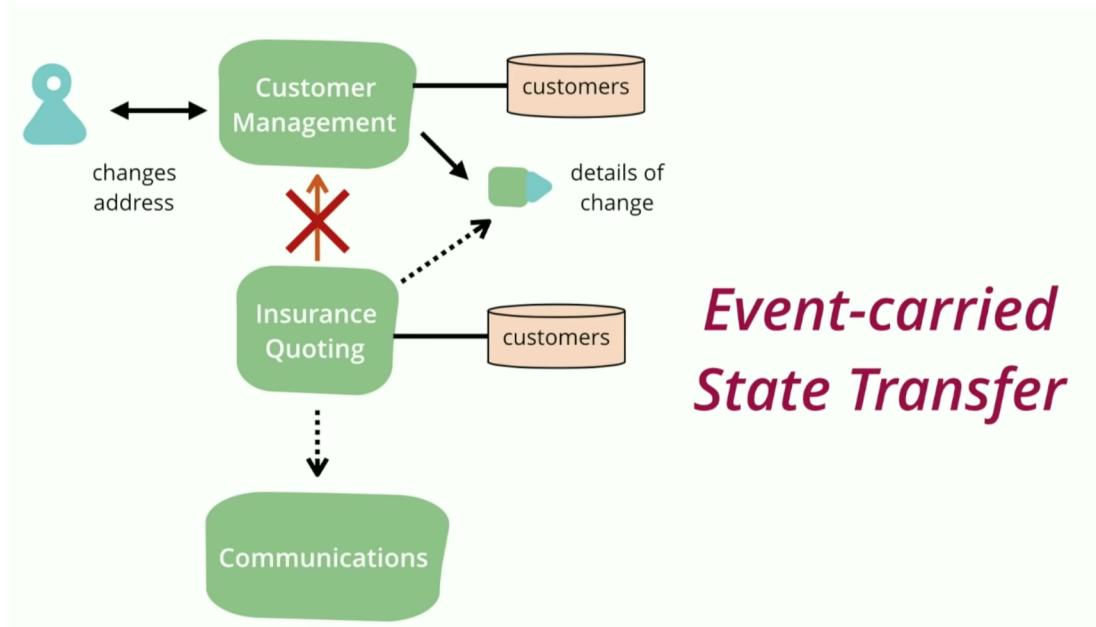




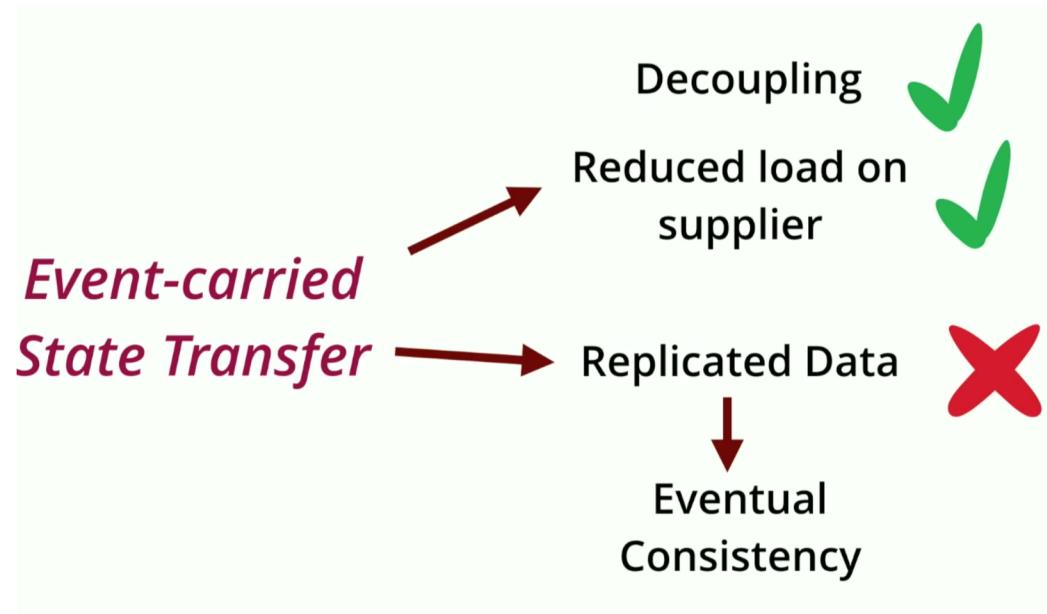
# Event-Carried State Transfer

- Upon a state change, data is transferred to other systems through events
- Greater decoupling as systems are not reliant on each other
- Lower latency, lower load on underlying systems

# Event-carried State Transfer



# Event-carried State Transfer





# Event Sourcing

- Whenever a change is made it is logged as an event
- System state can be rebuilt at any time from the event layer
- Large amounts of storage required at the kafka layer, doesn't work for use cases where a MV is required



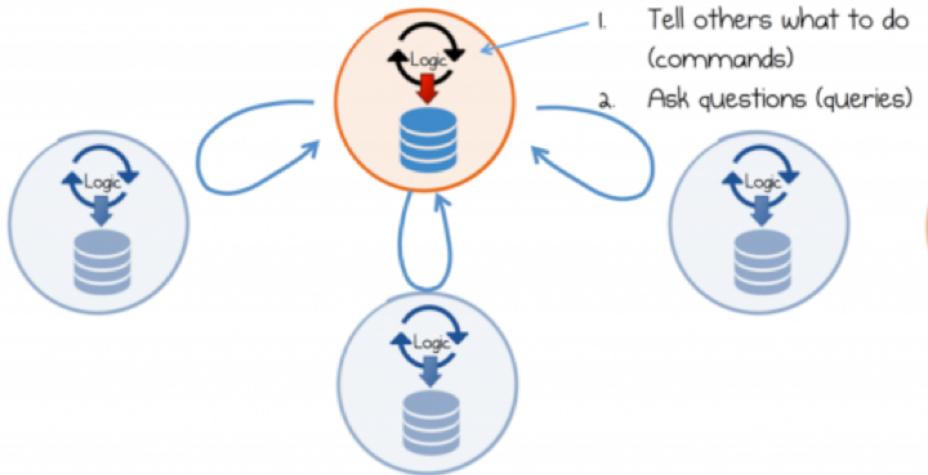
# CQRS

- Separate read and write
- Every method is either a command or a query
- Kafka can be used as a middle layer (write to kafka, pull the data out and write to a db to be read)

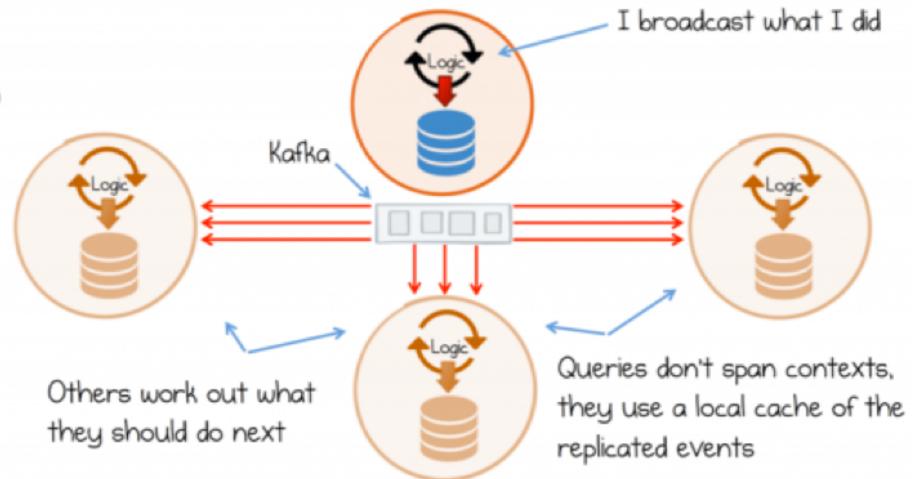
# A visual summary of commands, events, and queries



## REQUEST DRIVEN WAY



## EVENT DRIVEN WAY





# Introduction to Kafka

# Apache Kafka®: Open Source Streaming Platform

## Battle-Tested at Scale



The birthplace of Apache Kafka



More than 1 petabyte of data in Kafka



Over 4.5 trillion messages per day



60,000+ data streams



Source of all data warehouse & Hadoop data



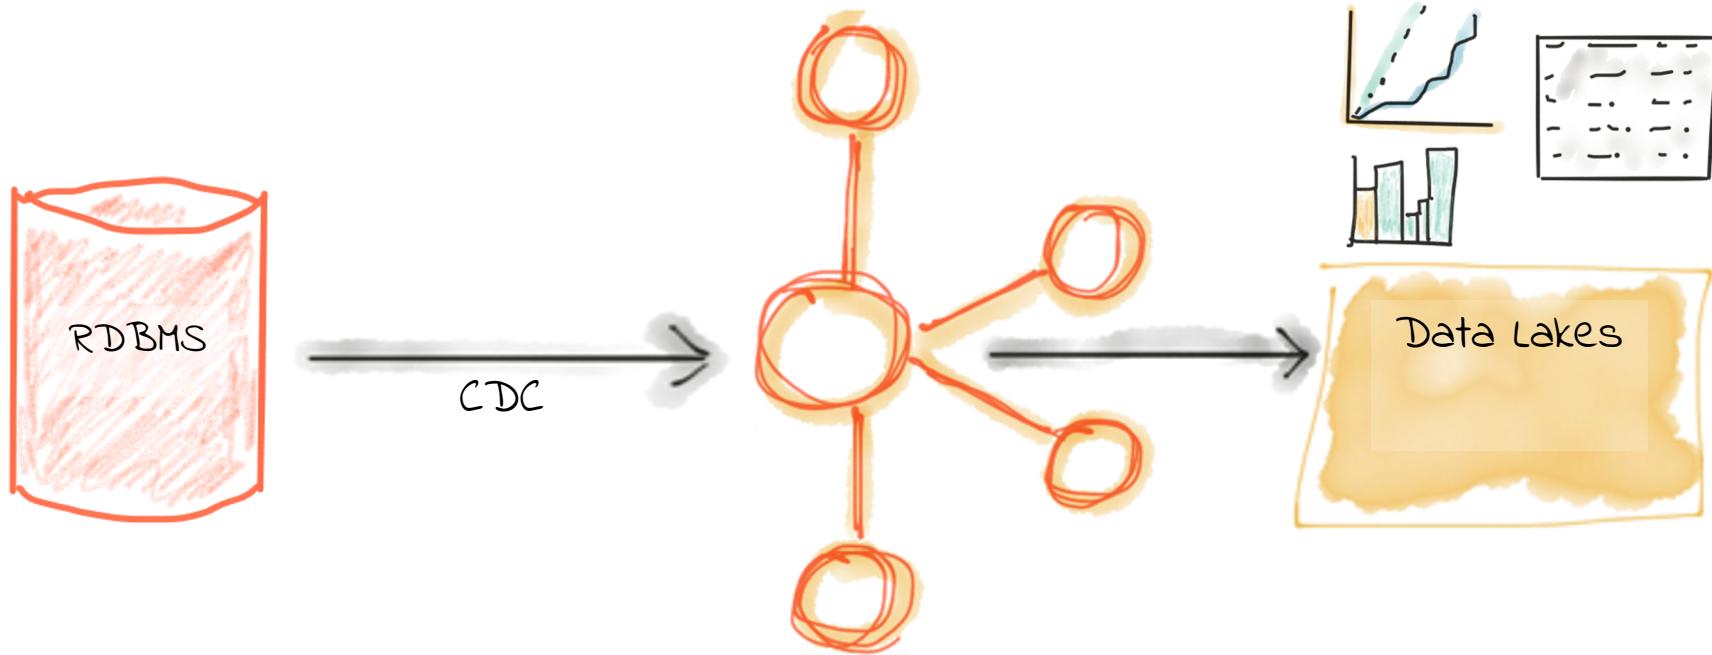
Over 300 billion user-related events per day



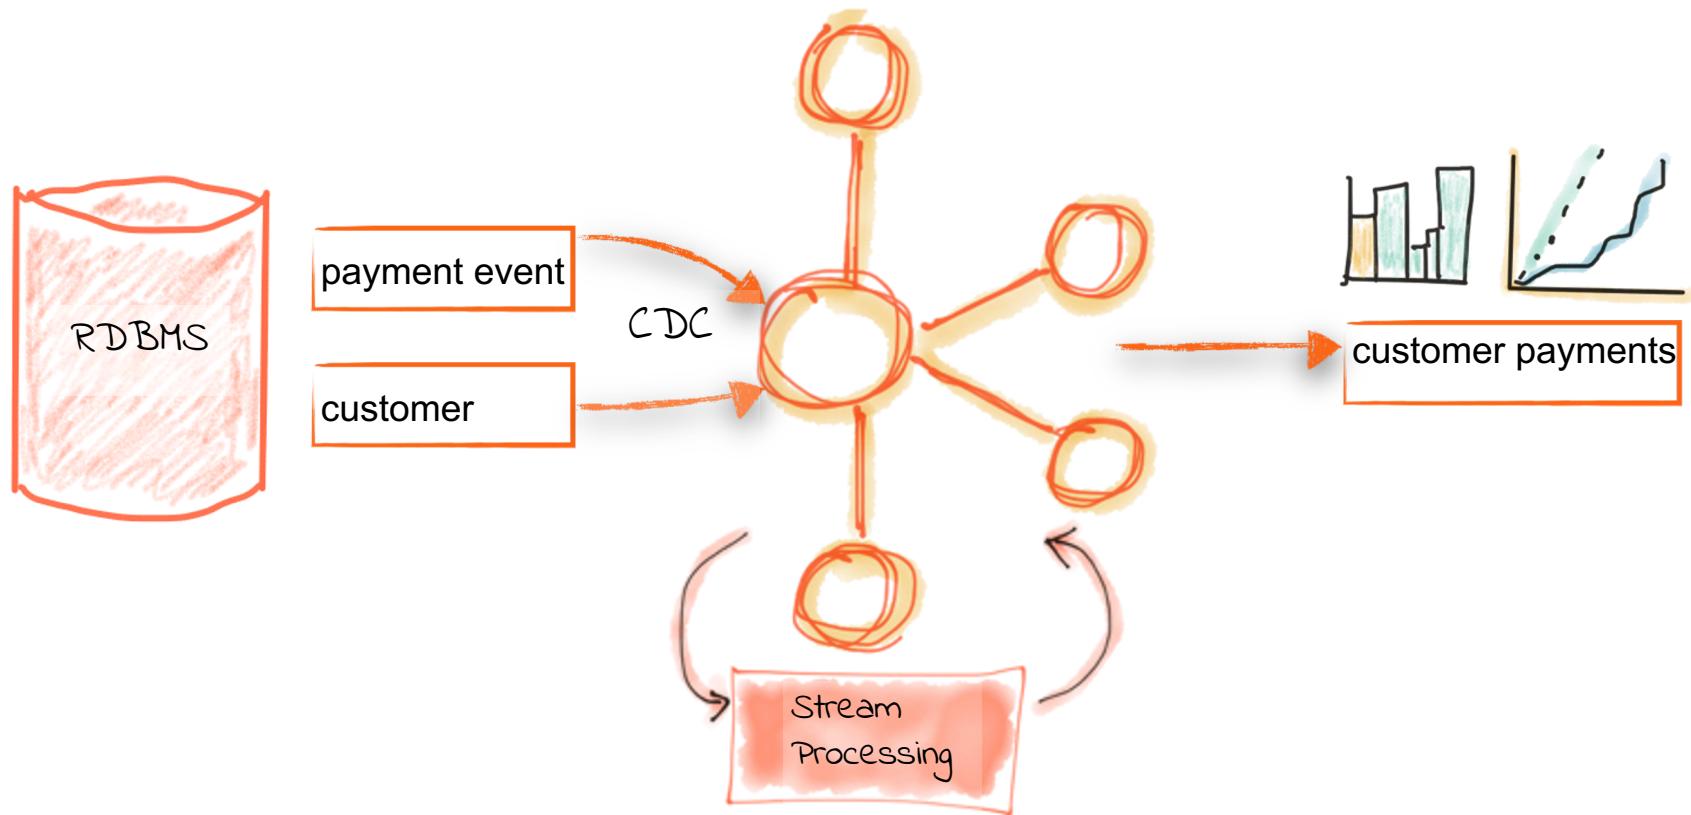
“

# Architectural Patterns with Apache Kafka

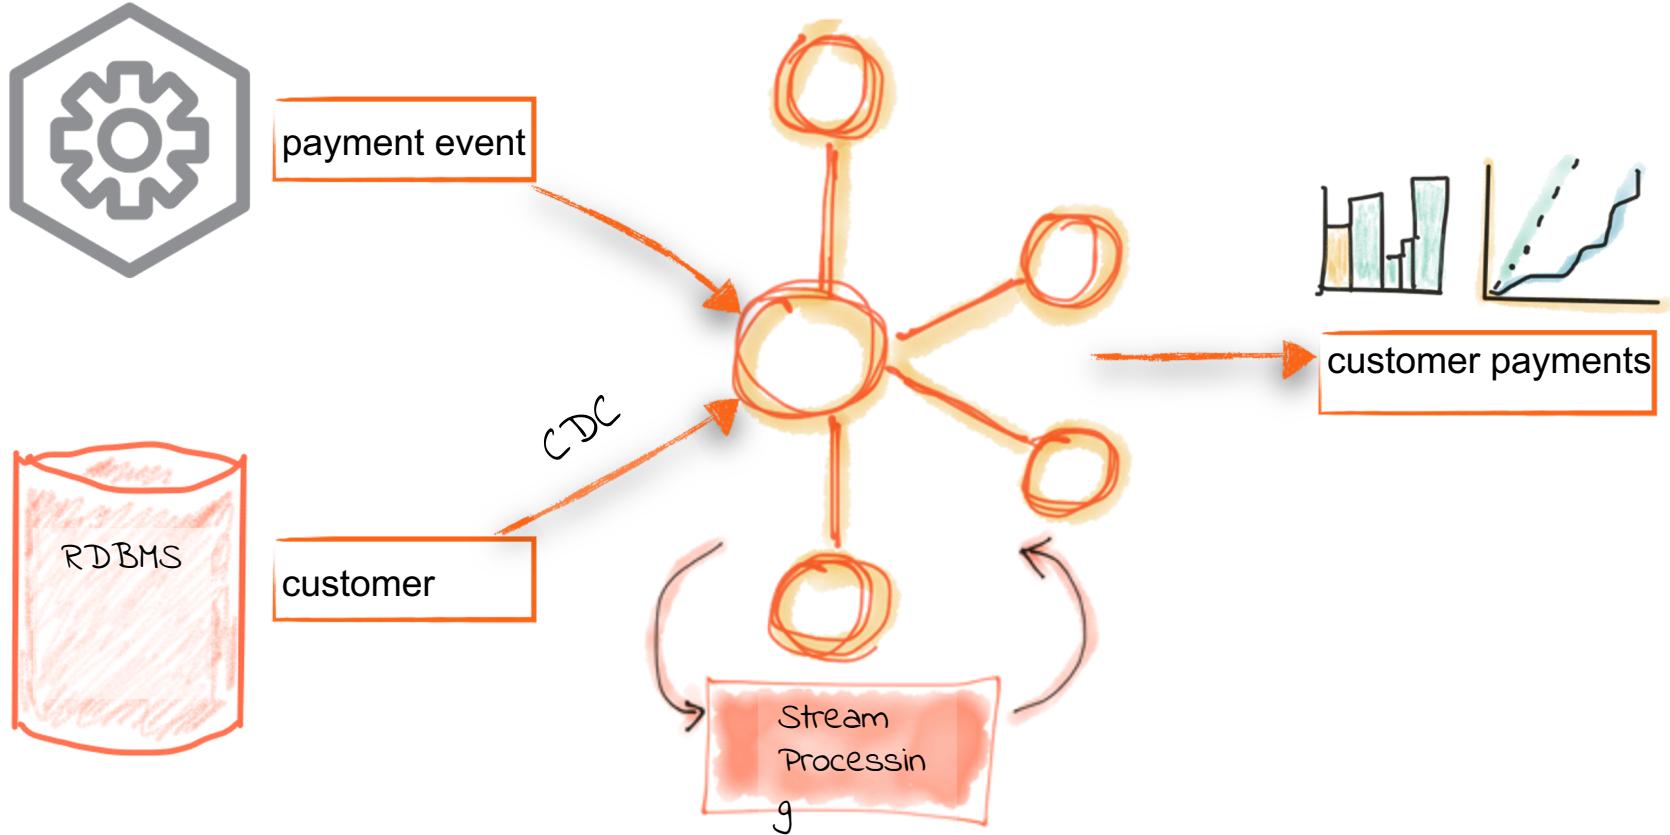
# Analytics - Database Offload



# Stream Processing with Apache Kafka and KSQL

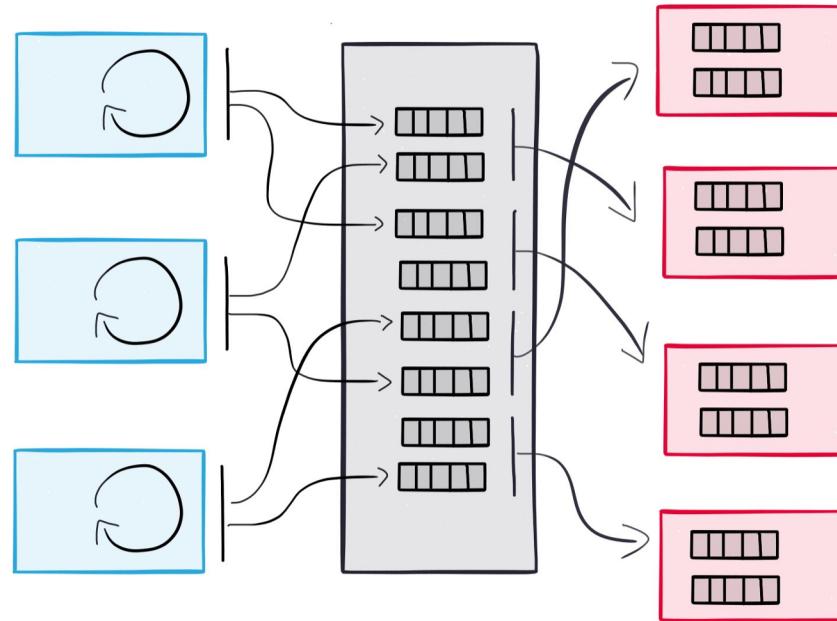


# Transform Once, Use Many





# Writers      Kafka cluster      Readers



# Scalability of a filesystem

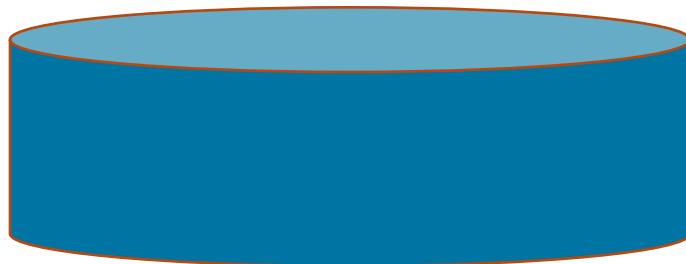


- hundreds of MB/s throughput
- many TB per server
- commodity hardware

# Guarantees of a Database

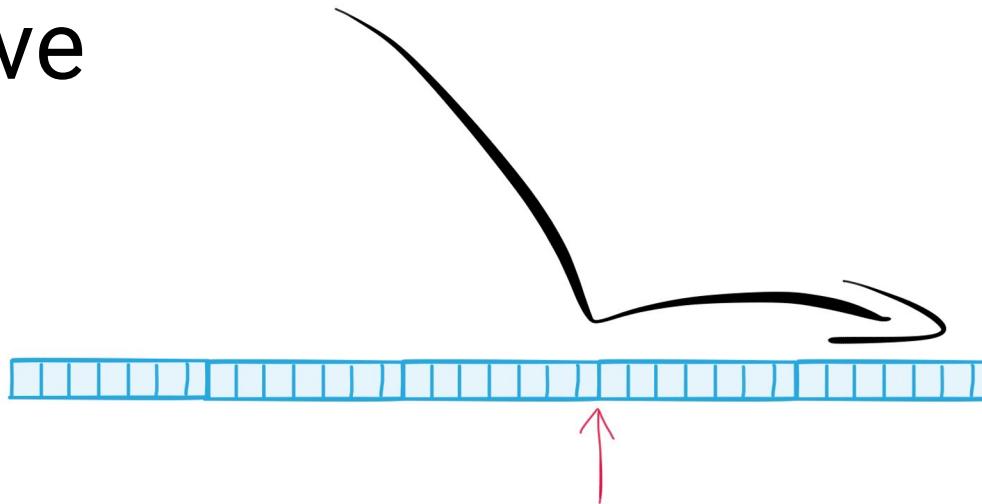


- Strict ordering
- Persistence





# Reset to any point in the shared narrative



Rewind & Replay

# Distributed by design



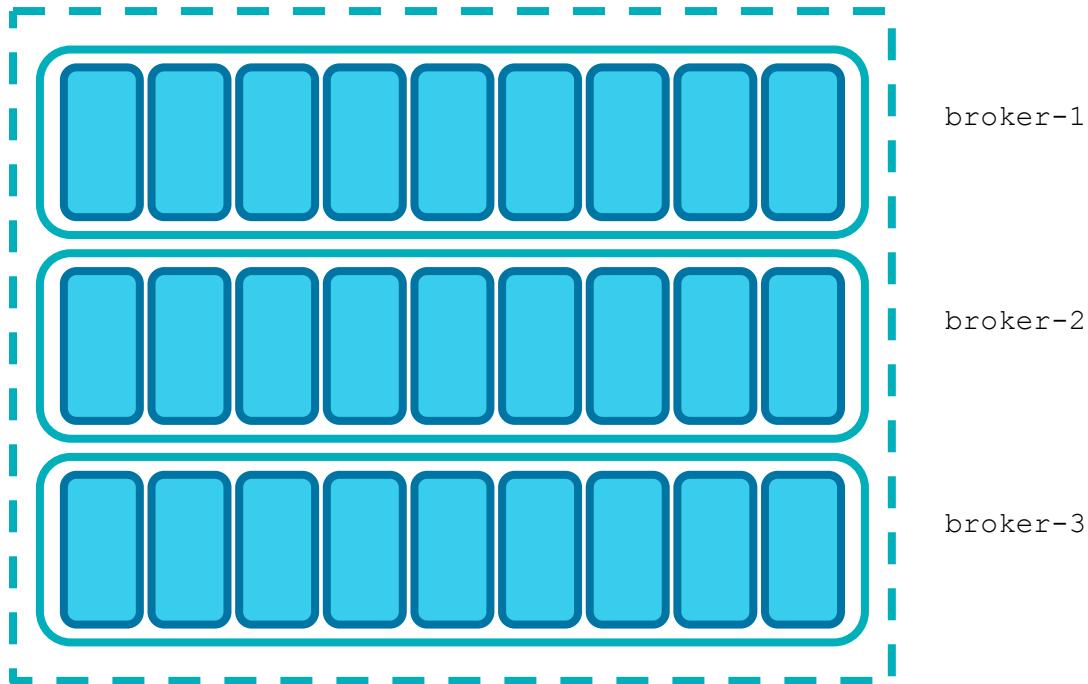
- Replication
- Fault Tolerance
- Partitioning
- Elastic Scaling

# Kafka Topics



my-topic

my-topic-partition-0



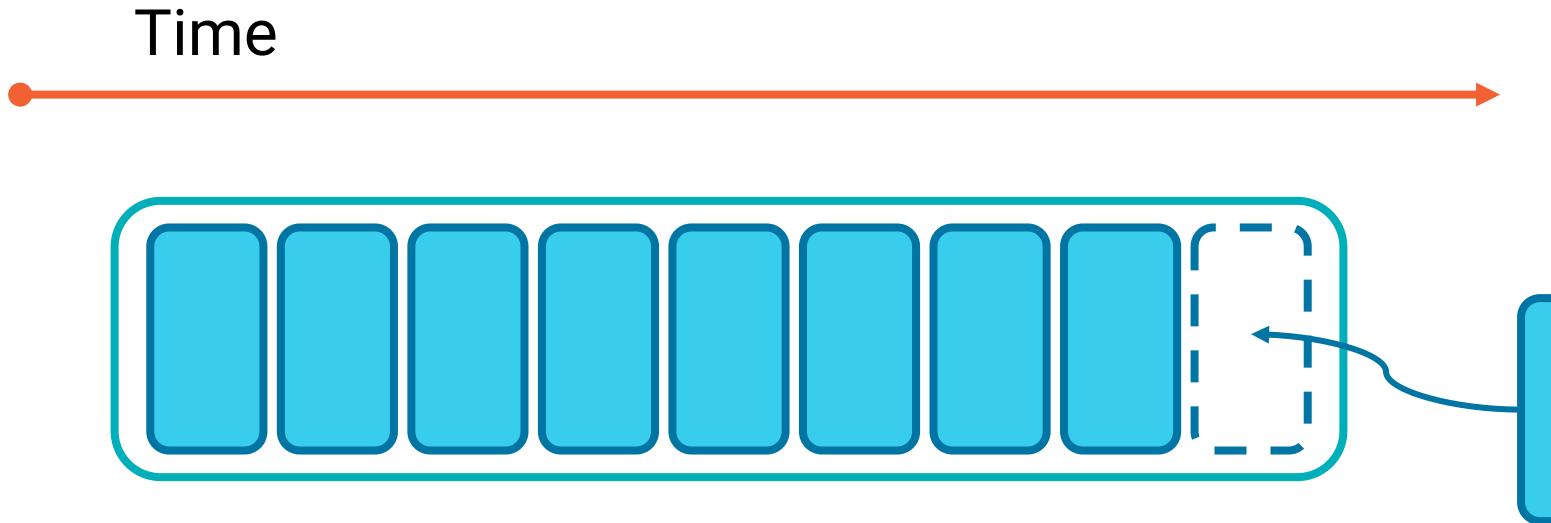
my-topic-partition-2

broker-1

broker-2

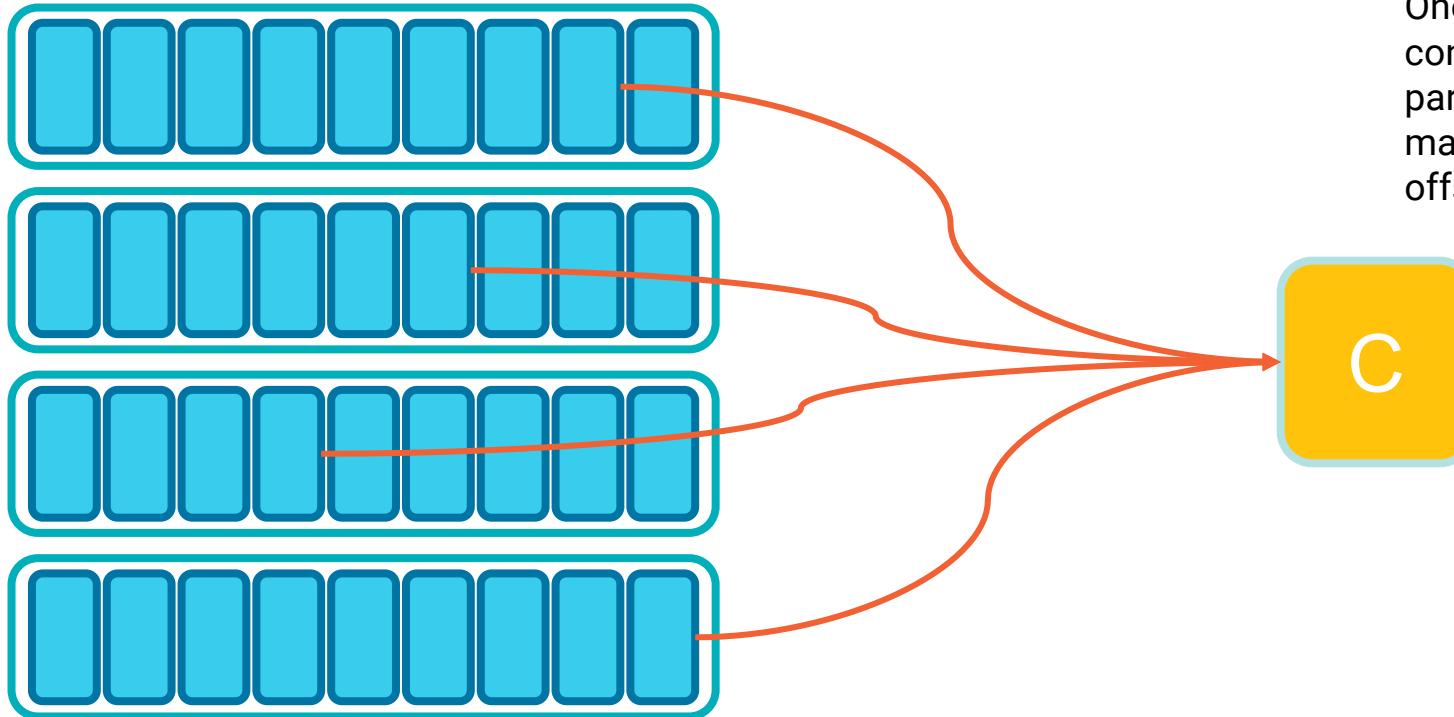
broker-3

# Producing to Kafka



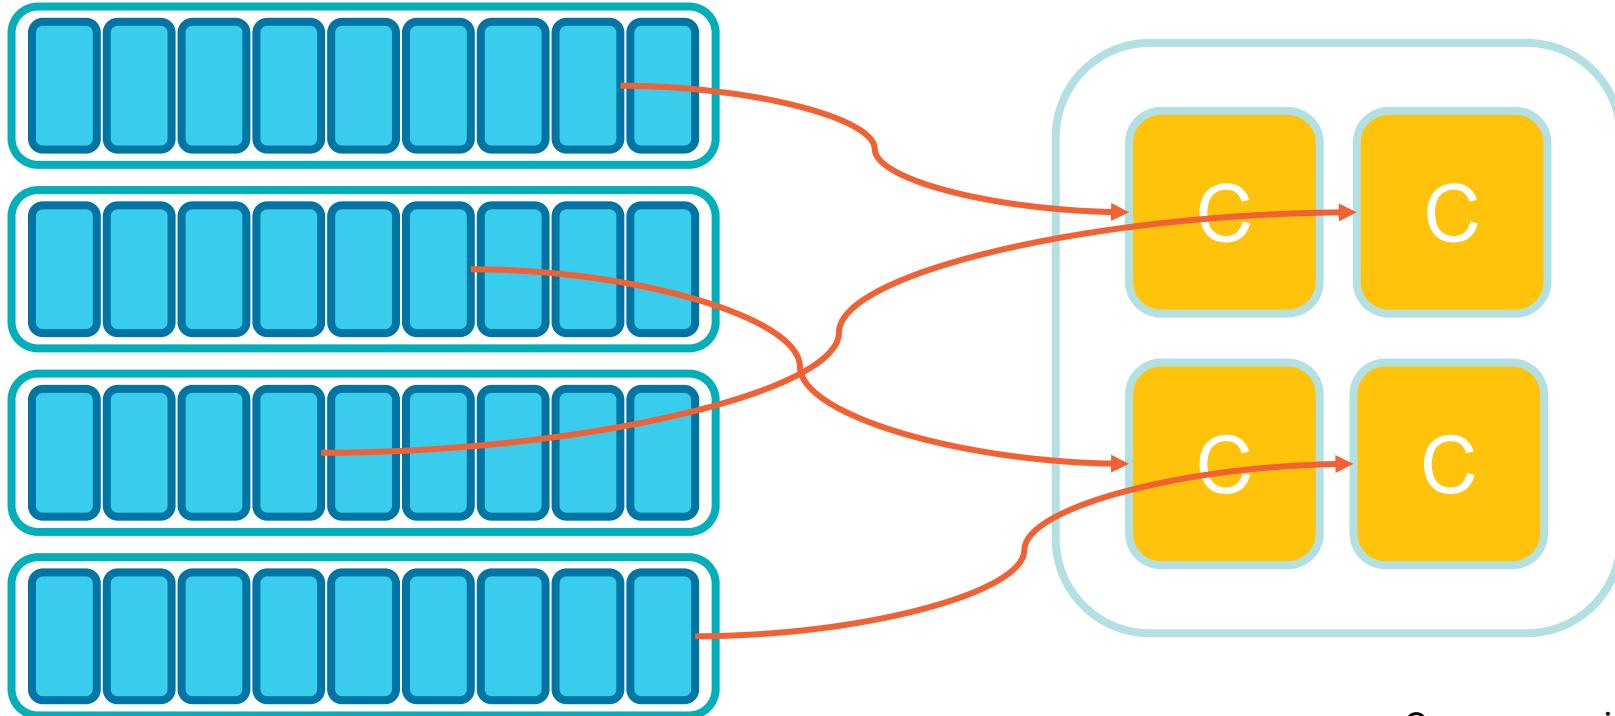


# Consuming From Kafka - Single Consumer



One consumer will consume from all partitions, maintaining partition offsets

# Consuming From Kafka - Grouped Consumers

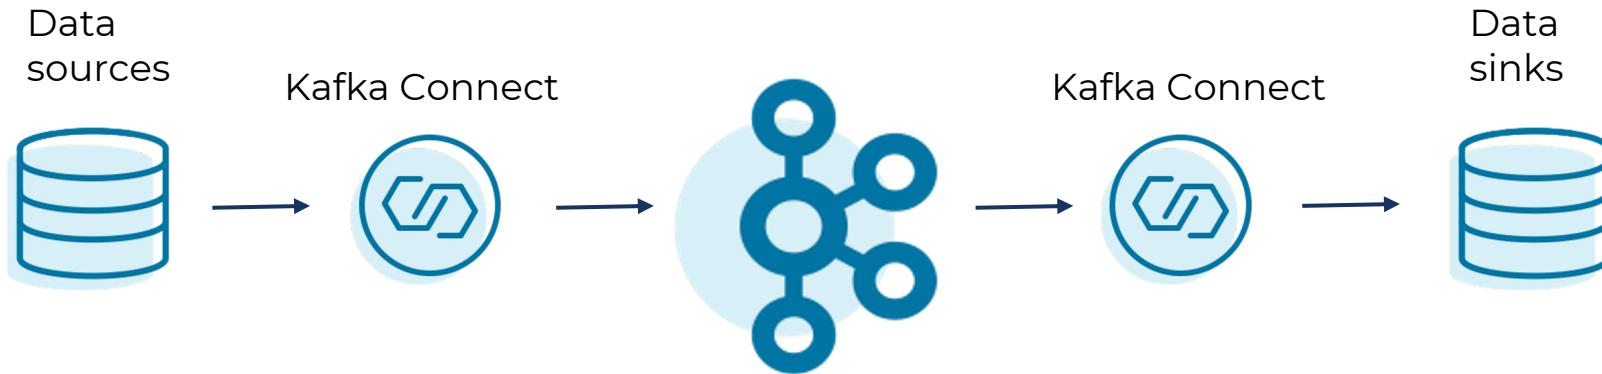


Consumers in a consumer group share the workload

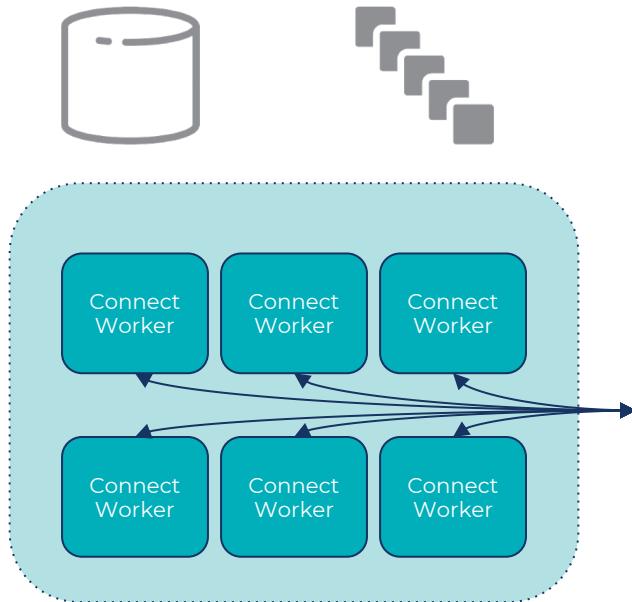
# Kafka Connect



No-Code way of connecting known systems (databases, object storage, queues, etc) to Apache Kafka



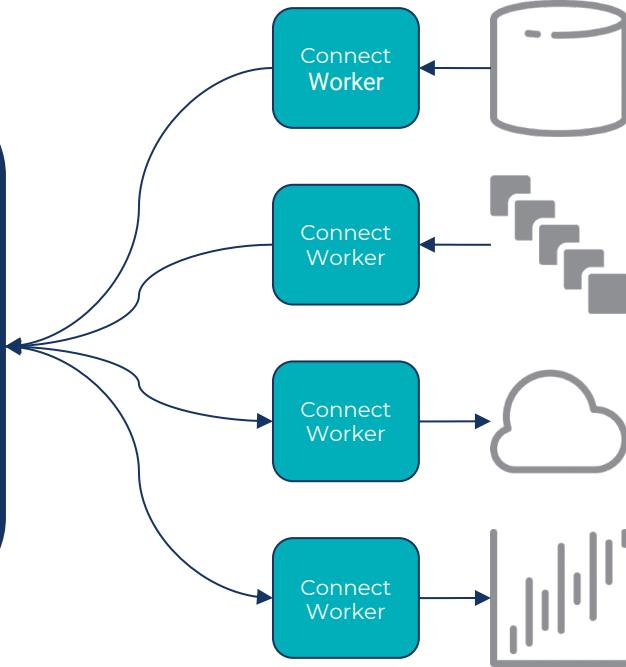
**Workers:** The JVM(s) that run **connectors** and **tasks**. Can be run in either **standalone** or **distributed** mode.



Connect Distributed Cluster



Connect Standalone Workers



# When to use Kafka's Streams API



- Mainstream Application Development
- To build core business applications
- Microservices
- Fast Data apps for small and big data
- Reactive applications
- Continuous queries and transformations
- Event-triggered processes
- The “T” in ETL
- <and more>

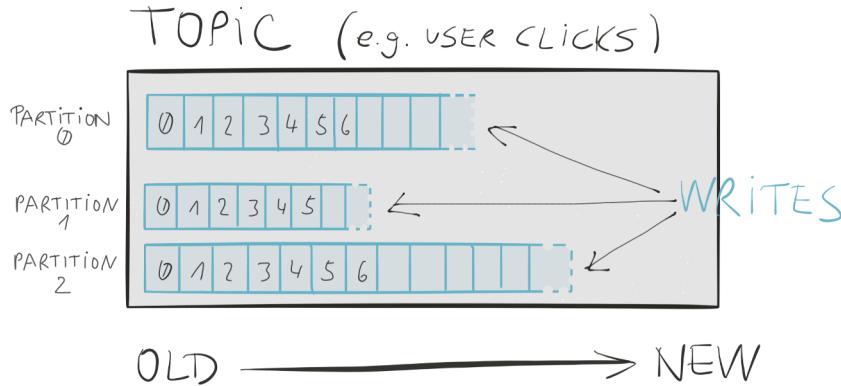
## Use case examples

- Real-time monitoring and intelligence
- Customer 360-degree view
- Fraud detection
- Location-based marketing
- Fleet management
- <and more>

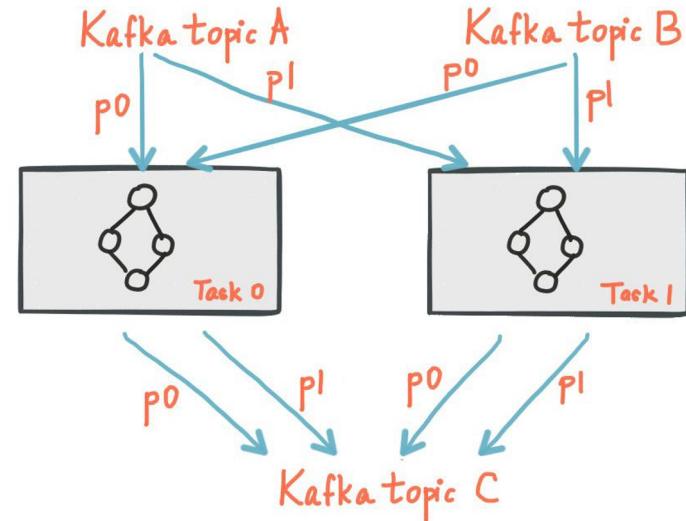
# Key concepts



## Kafka Core



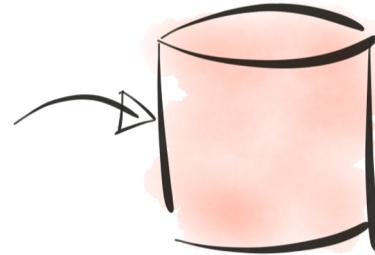
## Kafka Streams



STREAM → TABLE

CONTINUOUSLY  
UPDATING

0	1	2	3	4	5	6	7	8
---	---	---	---	---	---	---	---	---



EXAMPLE:  
CUSTOMER 360°

MATERIALIZED  
VIEW



# KSQL

```
CREATE STREAM fraudulent_payments AS  
SELECT * FROM payments  
WHERE fraudProbability > 0.8;
```

# VS.



```
object FraudFilteringApplication extends App {  
  
    val config = new java.util.Properties  
    config.put(StreamsConfig.APPLICATION_ID_CONFIG, "fraud-filtering")  
    config.put(StreamsConfig.BOOTSTRAP_SERVERS_CONFIG, "kafka-broker")  
  
    val builder: StreamsBuilder = new StreamsBuilder()  
    val fraudulentPayments: KStream[String, Payment] = builder  
        .stream[String, Payment]("payments-kafka-topic")  
        .filter(_ ,payment) => payment.fraudProbability > 0.8)  
  
    val streams: KafkaStreams = new KafkaStreams(builder.build(), config)  
    streams.start()  
}
```



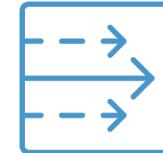
# KSQL example use cases



Data exploration



Data enrichment



Streaming ETL



Filter, cleanse, mask



Real-time monitoring



Anomaly detection



# Who should manage your Kafka cluster?

## Self Managed

- Granular configuration control
- Higher security control
- High DevOps burden

### Providers

- OS Kafka
- Confluent Community / Platform

## Fully Managed

- Faster time to live
- Lower operational burden
- Maximize uptime
- Lower operational control

### Providers

- MSK
- Confluent Cloud
- Heroku



CONFLUENT

# Resources



I Heart Logs: <https://www.confluent.io/ebook/i-heart-logs-event-data-stream-processing-and-data-integration/>

Designing Event Driven Systems: <https://www.confluent.io/designing-event-driven-systems/>

Kafka: The Definitive Guide: <https://www.confluent.io/resources/kafka-the-definitive-guide/>

Confluent Blog: <https://www.confluent.io/blog>

Reference Architecture: <https://www.confluent.io/resources/apache-kafka-confluent-enterprise-reference-architecture/>

Kafka Quickstart: <https://kafka.apache.org/quickstart>

Confluent Cloud Quickstart: <https://docs.confluent.io/current/cloud/quickstart.html>

Thank You!

agilbert@confluent.io

